

BLF as published — one forecast date

$b_0 : p = 0.5$

fresh belief every question



loop $t = 1 \dots T_{\max} = 10$

$(a_t, b_t) \leftarrow \text{LLM}(m)$ — one generation
search $\rightarrow$ LeakFilter $\rightarrow$ append, or submit



$K = 5$ trials $\rightarrow$ shrunken logit mean

$\hat{p} = \sigma(\alpha \cdot \text{mean logit } p_k)$, clamp $[0.05, 0.95]$



hierarchical Platt calibration

per-source intercept offsets

As run here — belief carried across dates

$b_0 = \text{yesterday's belief JSON}$

sequential; $p = 0.5$ in the parallel condition



loop $t = 1 \dots T_{\max} = 10$ — unchanged

LeakFilter replaced by the SQL date gate
(a stronger guarantee, not a weaker one)



$K = 3$ trials $\rightarrow$ same aggregator

$\alpha = 1/(1 + \text{var}(\text{logits}))$ — our instantiation



calibration omitted

no resolved training questions in the window