

BLF as published - one forecast date

b0 : p = 0.5

a fresh belief for every question



loop t = 1..Tmax = 10

one generation returns BOTH an action
and the updated belief



K = 5 trials -> shrunken logit mean

clamp to [0.05, 0.95]



hierarchical Platt calibration

per-source intercept offsets

As run here - belief carried across dates

b0 = yesterday's belief JSON

sequential; p = 0.5 in the parallel arm



loop t = 1..Tmax = 10 - unchanged

LeakFilter replaced by the SQL date gate,
a strictly stronger guarantee



K = 5 trials - same aggregator

alpha = $1/(1 + \text{var}(\text{logits}))$,
our instantiation of their alpha



calibration omitted

needs resolved training questions;
the window contains none