

FutureSim basicAgent as published

environment drives the clock

the agent calls `next_day()` when
it decides to advance

session: one tool call per turn

`query_df`, `search_news`, memory CRUD,
`submit_forecasts`, `next_day`

may end a day WITHOUT submitting

the standing forecast carries over

context cleared between sessions

structured memory and past
predictions are all that survive

hybrid semantic + keyword retrieval

LanceDB over 7.36M articles

As run here - inside our protocol

WE drive the clock

one session per scheduled date,
so all four harnesses share a calendar

session - unchanged

same action set; `query_df` is a real
pandas sandbox over the question table

MUST submit every date

their declining-to-update behaviour is
disabled to keep the tiers comparable

context cleared - unchanged

CRUD memory plus a separate
end-of-session memory phase

BM25 over the gated index

the same retrieval every harness gets,
so it is a constant not a confound